

The Information-Theoretic Essence of Intelligence: Compression, Verification, and the Limits of Self-Improvement

Anonymous

October 1, 2026

Abstract

We propose a *unifying interpretive framework* for understanding intelligence as an information-theoretic process: *intelligence is compression under verification*. We do not claim a new theory but rather a coherent synthesis that connects existing phenomena into a falsifiable worldview. Learning compresses experience into parameters, reducing the KL divergence between model and true distributions; verification provides an external lower bound on this divergence, grounding the compressed representation in reality. Without verification, compression converges to self-reported activity—the empty-loop fixed point (Paper 1). With isomorphic verification, structural blind spots emerge—the numerical precision blind spot (Paper 2). Both failure modes are corollaries of a single principle: *intelligence requires heterogeneous grounding*. We sharpen this into a *multiplicative claim*: effective intelligence factors as $I_{\text{eff}} = C \times V$, where C is compression capability and V is the accuracy-weighted coverage of verification; the empty loop is the $V=0$ degenerate point and the blind spot the V -is-correlated-with- C degenerate point. Recent work showing truth to emerge as a compression artifact stops one step short: emergent truth remains an *unreconciled loan* until verified. We further report an engineering corollary—*verification as reflex*: verification becomes economical only when its own cost is amortized by compression, which we implement as a three-tier architecture (reflex/memory/evolution) with production readings. We formalize the *amortization rate* as an information compression ratio $R = I_{\text{serial}}/I_{\text{parallel}}$, showing that scaling laws describe the prepayment of dimensions and energy for exponential search space compression, bounded by thermodynamic constraints. The framework unifies the manifold hypothesis, amortized inference, symbol grounding, and Goodhart’s law into a coherent worldview. Its value lies not in novelty but in *unification and falsifiability*: it explains existing phenomena, predicts failure modes validated in companion papers, and provides a common language for discussing intelligent systems.

1 Introduction: The Missing Foundation

The past decade has produced extraordinary empirical progress in machine intelligence—scaling laws, self-improving agents, LLM-as-judge evaluation—without a corresponding theoretical foundation. We observe:

- **Scaling laws** [1, 2]: performance improves as a power law of scale, but *why*?
- **Self-improvement** [3, 4]: agents improve from experience, but *when does this fail*?
- **LLM-as-judge** [11]: models evaluate models, but *what are the limits*?

We propose that these phenomena share a common foundation: *intelligence is an information-theoretic process of compression under verification*. This is not a new theory but a *unifying interpretive framework*—a coherent synthesis that connects existing phenomena into a falsifiable worldview. Its value lies in:

1. **Unification:** Connecting scaling laws, self-improvement failure modes, and evaluation limits into a single framework;
2. **Falsifiability:** Making testable predictions (Empty-Loop Theorem, Blind Spot Theorem) that are validated in companion papers;
3. **Explanation:** Providing a common language for discussing why intelligent systems succeed or fail;
4. **Guidance:** Offering practical principles for building reliable intelligent systems.

We explicitly position this work as an *interpretive framework* rather than a new theory. The components—compression, verification, amortization—are well-established in their respective fields. Our contribution is their synthesis into a coherent whole with falsifiable predictions.

2 The Core Framework

2.1 Intelligence as Constrained Optimization

Definition 1 (Intelligence, operational). *Intelligence is the process of maximizing compression subject to verification constraints:*

$$\text{Intelligence} = \max_{\theta} \text{Compression}(\theta) \quad \text{subject to} \quad \text{Verification}(\theta) \geq \varepsilon \quad (1)$$

where:

- $\text{Compression}(\theta) = -D_{KL}(p_{\text{true}} \| p_{\text{model}_{\theta}})$ (compression rate, larger is better)
- $\text{Verification}(\theta) = \text{external lower bound provided by verification (constraint)}$
- $\varepsilon = \text{minimum acceptable verification level}$

Compression alone is insufficient. Without external verification, the model cannot estimate D_{KL} —its self-estimate is part of the compressed representation. Verification provides an *external lower bound* on D_{KL} .

2.2 The Empty-Loop Theorem

Proposition 1 (Empty-Loop Theorem, formal). *Let π_t be the policy at time t , and let $V : \mathcal{E} \rightarrow \{0, 1\}$ be the verification function. If $V(e) = 0$ for all $e \in E$, then:*

$$\lim_{t \rightarrow \infty} D_{KL}(p_{\text{true}} \| \pi_t) = D_{KL}(p_{\text{true}} \| \pi_0) \quad (2)$$

Moreover, the system converges to a self-reported activity fixed point: internal metrics grow without bound while D_{KL} remains constant.

Derivation:

1. Learning aims to reduce $D_{KL}(p_{\text{true}} \| p_{\text{model}})$.
2. Without verification, the model estimates D_{KL} using its own distribution.
3. This estimate is itself part of D_{KL} (self-reference).

4. Therefore, the model cannot know how far it is from p_{true} .
5. The system converges to a fixed point where it “appears to learn” without actual progress.

Empirical validation: Paper 1 demonstrates empty-loop convergence in controlled experiments (PoI-Gate-1) and a 50,000-cycle production failure.

2.3 The Blind Spot Theorem

When verification is *isomorphic* to the generator (same representation space), structural blind spots emerge:

Proposition 2 (Blind Spot Theorem, formal). *Let $U \subset \mathbb{R}^d$ be the judge’s discrimination space and $W \subset \mathbb{R}^d$ the error space. The judge’s pass rate is:*

$$\rho = 1 - \frac{\|P_U W\|_F}{\|W\|_F} \quad (3)$$

where P_U is the projection onto U and $\|\cdot\|_F$ is the Frobenius norm. For isomorphic judges ($U_{\text{judge}} \approx U_{\text{generator}}$), the blind spot is maximal.

Derivation:

1. The judge’s discrimination direction U is determined by training data.
2. Isomorphic judges share training data with the generator.
3. Therefore, $U_{\text{judge}} \approx U_{\text{generator}}$.
4. The judge is blind to error directions W orthogonal to U .
5. Blind spot size is determined by overlap(U, W).

Empirical validation: Paper 2 demonstrates the blind spot in high-precision numerical perturbations (pass rate 37–48% for $\pm 0.1\%$ perturbations).

3 Verification as Reflex

The framework above treats verification as a constraint. Its *economics* decides whether the constraint is ever enforced: verification that costs as much as the task it guards is never invoked. We therefore require verification to be cheap enough to become a *reflex*—executed on every write, without deliberation.

3.1 The Three-Tier Architecture

- **L0 (reflex):** a verifier embedded in write paths (memory-ingest gate, training-delivery gate, outbound-receipt gate) with latency budgets at the tens-of-milliseconds scale. Production reading: median 30.1 ms, p95 42.2 ms, throughput 118 items/s.
- **L1 (memory):** each verdict—including failures and errata—feeds a corpus that periodically retrains the verifier; verification experience is itself compressed into 25.7 MB of adapter weights.

- **L2 (evolution)**: corpus sampling respects a three-level truth hierarchy (T1 sensor-grade, T2 human adjudication, T3 mechanical rules), and per-domain adapters share one backbone. An anti-self-reference guard bounds what may be self-trained: only small verifiers, rewarded only against recomputable ground truth.

3.2 The Cost Curve

Let C_{first} be the marginal cost of the first occurrence of an error class (all human and compute cost of discovering it), and C_{second} the marginal cost of the second occurrence after the reflex has been trained on the first. The architecture drives $C_{\text{second}} \rightarrow 0$: *first mistakes are expensive; second mistakes are free to prevent*. Six days of organizational logs across eight error classes exhibit this curve (Appendix A); we propose the curve itself, not any single point, as the correct economic statement of “intelligence is compression under verification.”

4 Borrowed Time–Space

Training compresses temporal experience (the input–output process) and situational structure into high-dimensional geometry; inference *borrow*s back time (prediction) and space (generalization). The borrowed structure is a lossy copy, not the original. Verification is the audit that reconciles the loan.

A formal route runs through the rate–distortion–perception (RDP) tradeoff: given rate R and distortion D , perceptual quality P trades off against both. We extend to a four-way tradeoff **RDP-V**: every unit of verification coverage V must be paid for in rate or tolerated in distortion. To our knowledge the verification axis of this tradeoff is unoccupied in the literature (survey, October 2026); we claim it as the formal seat of the multiplier $I_{\text{eff}} = C \times V$.

5 The Amortization Rate

5.1 Information-Theoretic Definition

Definition 2 (Amortization Rate, information-theoretic). *The amortization rate is:*

$$R = \frac{I_{\text{serial}}}{I_{\text{parallel}}} = \frac{L \cdot \log_2(n)}{d^2 \cdot \log_2(\text{precision})} \quad (4)$$

where L is sequence length, n vocabulary size, d model dimension, and *precision* is the bit-width of parameters.

R measures the degree to which exponential sequence space information is compressed into parameter space.

5.2 Derivation from Bayesian Inference

1. **Bayesian inference**: $p(y|x) = \int p(y|x, \theta)p(\theta)d\theta$ (exponentially complex).
2. **Amortized inference**: Learn approximate posterior $q_\phi(\theta|\text{data})$; inference becomes $p(y|x) \approx \int p(y|x, \theta)q_\phi(\theta|\text{data})d\theta$ (single forward pass).
3. **Scaling laws**: Larger models (more parameters) yield better approximations q_ϕ , reducing $D_{\text{KL}}(p_{\text{true}}||p_{\text{model}})$.

5.3 Thermodynamic Constraints

Landauer’s principle [14] states that erasing one bit requires energy $kT \ln 2$. Each forward pass “erases” uncertainty about the next token:

$$E_{\text{forward}} \geq kT \ln 2 \cdot (d^2 \cdot L) \cdot \log_2(\text{precision}) \quad (5)$$

As $d \rightarrow \infty$, $R \rightarrow \infty$ but $E_{\text{forward}} \rightarrow \infty$ faster. There exists an optimal scale d^* maximizing R/E_{forward} .

Testable prediction: The performance-per-energy ratio peaks at some d^* and declines thereafter.

5.4 Production Readings

The amortization machinery is not hypothetical. An organizational deployment supplies: (i) a 2,053-row decision dataset converted from the full organizational mailbox, including a 107-judgment user-verdict aligned subset (T2 truth); (ii) a three-state verdict classifier trained on 1,454 labeled verdicts under a frozen split (test accuracy 88.51%, ECE 0.072, +22.3 points over the isomorphic baseline of 66.2%); (iii) an external protocol measurement over 40 seeded synthetic CVs (92.6% generator-agreement, Brier 0.069, ECE 0.138, with 2/27 voluntary abstentions), with the answer key committed and timestamped before any model call.

6 Unified Worldview

The framework yields a coherent worldview:

Intelligence is compression under verification.

Learning compresses experience (reducing D_{KL}); verification grounds the compression (providing external lower bounds). Without verification, compression converges to empty loops (Paper 1). With isomorphic verification, blind spots emerge (Paper 2). The amortization rate quantifies the computational structure of this process, bounded by thermodynamics.

6.1 Falsifiable Predictions

The framework makes testable predictions:

1. **Empty-loop convergence:** Unverified self-improving agents exhibit $L = 0$ despite activity (Paper 1, validated).
2. **Blind spot sharing:** Isomorphic judges exhibit correlated pass rates (Paper 2, validated).
3. **Δ as predictor:** The discrimination signal Δ predicts pass rates (Paper 2, validated).
4. **Amortization rate scaling:** R grows as a power law of model scale (future work).
5. **Thermodynamic bound:** Performance-per-energy peaks at some d^* (future work).
6. **Reflex threshold:** embedding L0 reflex verifiers reduces same-class error recurrence below 5% within one week (organizational baseline: four of five error classes recurred at least twice pre-reflex).

7. **Multiplicative degradation:** verifiers trained on the same distribution as the system they verify gain less per unit of nominal coverage than heterogeneous verifiers—retrospectively testable on the three blind-spot families of Paper 2.
8. **T2 sparsity—value inversion:** user-adjudication signals are rarer than 5% of all verdicts yet contribute more marginal retraining value than any T3 source—testable directly on the decision dataset.

7 Related Work

7.1 Scaling Laws

Kaplan et al. [1] and Hoffmann et al. [2] establish empirical scaling laws. We provide the theoretical foundation: scaling is the amortization of search space compression.

7.2 Self-Improvement

Voyager [3] and Reflexion [4] demonstrate self-improvement with verification. TRT [5] and Crescent [6] claim self-improvement without external feedback, but their verification is internal: TRT uses self-generated verification signals at test time; Crescent uses majority voting (internal consistency) in mathematical reasoning. We distinguish: *internal verification in limited settings* (test-time, single-task, unique-answer domains) vs. *long-term experience accumulation with external verification* (our focus). The Empty-Loop Theorem (Paper 1) applies to the latter. Crescent’s own citation of Kambhampati et al. [7]—questioning self-generated data training—is consistent with our formalization.

7.3 Memory Systems

Recent work on agent memory [8, 9, 10] focuses on admission control, retrieval, and temporal knowledge graphs, but assumes stored information is trustworthy. A-MAC [8] notes that “hallucinated or outdated information can propagate errors” but relies on LLM-based admission evaluation. Our framework provides the theoretical foundation: without external verification, memory systems are vulnerable to empty-loop convergence (Paper 1) and blind spot propagation (Paper 2).

7.4 LLM-as-Judge

MT-Bench [11] establishes LLM-as-judge. We identify structural limits: isomorphic judges share blind spots (Paper 2).

7.5 Symbol Grounding

Harnad [12] poses the symbol grounding problem. We provide the engineering solution: verification as grounding.

7.6 Goodhart’s Law

Goodhart’s Law [13]—“when a measure becomes a target, it ceases to be a good measure”—is related but distinct. Empty-loop convergence is not merely metric gaming; it is the *absence of any metric*, replaced by self-reported narrative.

8 Discussion

8.1 For Architecture Design

Maximize R while maintaining verification channels. Attention mechanisms achieve high R through geometric operations; verification requires heterogeneous information channels.

8.2 For Agent Design

Self-improving agents require external verification. The Verification Savings Rate (VSR, Paper 1) is a leading health indicator.

8.3 For Evaluation

LLM-as-judge requires Δ pre-deployment checks (Paper 2). Isomorphic judges are structurally insufficient.

8.4 Philosophical Implications

The framework provides an operational answer to the symbol grounding problem: verification is the ground truth interface that prevents self-referential drift. Without V , the system’s symbols refer only to other symbols, not to external reality.

9 Conclusion

We have proposed a unifying interpretive framework: intelligence is compression under verification. We do not claim a new theory but rather a coherent synthesis that connects existing phenomena into a falsifiable worldview. The framework explains scaling laws, predicts self-improvement failure modes (validated in Papers 1 and 2), formalizes evaluation limits, and unifies diverse phenomena. Its value lies in unification and falsifiability: it provides a common language for discussing intelligent systems and makes testable predictions about when they fail. The framework provides both theoretical foundation and practical guidance for building reliable intelligent systems.

References

- [1] Kaplan, J., et al. (2020). Scaling laws for neural language models. *arXiv:2001.08361*.
- [2] Hoffmann, J., et al. (2022). Training compute-optimal large language models. *arXiv:2203.15556*.
- [3] Wang, G., et al. (2023). Voyager: An open-ended embodied agent with large language models. *arXiv:2305.16291*.
- [4] Shinn, N., et al. (2023). Reflexion: Language agents with verbal reinforcement learning. *NeurIPS*.
- [5] Zhuang, Y., et al. (2026). Test-time recursive thinking. *arXiv:2602.03094*.
- [6] Sun, Z., et al. (2025). Crescent: Self-improvement without external seeds. *arXiv:2502.13441*.
- [7] Kambhampati, S., et al. (2024). Can large language models really improve by self-critiquing their own plans? *arXiv:2410.01465*.

- [8] Anonymous. (2026). A-MAC: Memory admission control for LLM agents. *arXiv:2603.04549*.
- [9] Packer, C., et al. (2023). MemGPT: Towards LLMs as operating systems. *arXiv:2310.08560*.
- [10] Rasmussen, D., et al. (2025). Zep: A temporal knowledge graph architecture for agent memory. *arXiv:2501.13956*.
- [11] Zheng, L., et al. (2023). Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. *NeurIPS*.
- [12] Harnad, S. (1990). The symbol grounding problem. *Physica D*, 42(1-3), 335–346.
- [13] Goodhart, C. A. E. (1975). Problems of monetary management: The U.K. experience. *Papers in Monetary Economics*.
- [14] Landauer, R. (1961). Irreversibility and heat generation in the computing process. *IBM Journal of Research and Development*, 5(3), 183–191.
- [15] Anonymous. (2026). Verification Learning: The Necessity of External Validation for Self-Improving Agents. *Companion paper (Paper 1)*.
- [16] Anonymous. (2026). The Numerical Precision Blind Spot of LLM-as-Judge. *Companion paper (Paper 2)*.

9.1 Abstaining Judges and Generative Verifiers (2025–2026)

Judges that selectively abstain [5] formalize what our three-state classifier runs in production (with calibration readings). Generative process reward models trained on thousands of labels [1] share our small-sample philosophy; our reflex layer is their production deployment. Reward-hacking analyses [2] are the RLVR form of the blind-spot theorem. On self-correction, the illusion that models correct others better than themselves [4] matches our open-book blind-spot family. On governance, the concept of *verified self-improvement* entered policy discourse in 2026 explicitly unfinished [3]; this paper reports an organizational-scale engineering answer, on the defensive side of the narrative.

References

- [1] Process reward models that think. *arXiv preprint arXiv:2504.16828*, 2025.
- [2] Llms gaming verifiers: Rlvr can lead to reward hacking. *arXiv preprint arXiv:2604.15149*, 2026.
- [3] Preparing for recursive self-improvement. OpenAI policy proposal, UN General Assembly week, 2026.
- [4] The self-correction illusion: Llms correct others but not themselves. *arXiv preprint arXiv:2606.05976*, 2026.
- [5] A. Kalai et al. Evaluating llms for accuracy incentivizes disaggregation. 2026.

A Organizational Lessons as Compression-Distortion Data

Eight error classes from six days of organizational logs (2026-09-26 to 10-01), each mapped to its distortion type (the compression-theoretic reading), the verification antibody deployed, and the cost-curve point it contributes (Section 3). All rows are traceable to organization-internal artifacts; identifiers are elided for operational security.

Error class	Distortion type	Verification antibody	Cost-curve point
Verifier ceiling (66.2%)	lossy summary drops discriminative structure	heterogeneous re- training (LoRA, three- state verbalizer; 88.51%)	$C_1=2$ GPU-h; $C_2 \rightarrow 0$
Self-report over- write	verdict field overwrit- ten by judge output (same-source)	artifact-anchored truth + dual-layer recompute (500 rows recalled)	near-miss; audit cost only
Open-book fea- tures	verifier features con- tain the judged output	feature-pipeline excision; re-run under frozen crite- ria	same C_1 , +22.3 pt
Equivalence mis- labeling	mechanical rule can- not canonicalize sur- face forms (13.8% con- flict)	rule demoted to sampled pool; criterion registry	3 verified instances
Fake training	200 steps logged, zero weight update (six checkpoints, identical MD5)	delivery telemetry: trainable > 0 , grad- norm > 0 , distinct checkpoints	single telemetry check
Synthetic- benchmark mirage	local 175 MB vs. production 4.9 GB (legacy stock unmod- eled)	test sets sampled from production distribution	two rollbacks, one fix
Theater metrics	income self-produced 100%, external settle- ment $B=0$ for two months	world-reply criterion (ex- ternal readings as final measure)	institutional antibody
Platform- boundary fric- tion	publish at- tempts die at captcha/permission/schem- edges	pre-registered protocol templates; boundary checklist	one manual step re- mains

Table 1: Eight organizational error classes as compression-distortion data. Each antibody, once trained into the reflex layer, drives the marginal cost of the second occurrence of its class toward zero—the economic statement of $I_{\text{eff}} = C \times V$.

Reading the table. The first column is what failed; the second is *why* the failure was invisible to the system that produced it (always a compression artifact: a dropped field, an overwritten bit, an unmodeled stock); the third is the heterogeneous grounding that caught it; the fourth is the amortization evidence. Four of the eight classes recurred at least twice before their antibodies were institutionalized—the baseline against which the reflex-threshold prediction (Section on falsifiable predictions) will be scored.